

# Machine Learning With TensorFlow

X433.7-001 (2 semester units in COMPSCI)

Instructor Alexander I. Iliev, Ph.D.

# Course Content Outline

- **Machine Learning With TensorFlow®**

- Introduction, Python - pros and cons
- Python modules, DL packages and scientific blocks
- Working with the shell, IPython and the editor
- Installing the environment with core packages
- Writing “Hello World”

HW1 (10pts)

- **Tensorflow and TensorBoard basics (cont.)**

- Ecosystem, Competition, Users
- Linear algebra recap
- Data types in Numpy and Tensorflow
- Basic operations in Tensorflow
- Graph models and structures with Tensorboard

- **TensorFlow operations**

- Sessions, graphs, variables, placeholders
- Overloaded operators
- Using Aliases

- **Data Mining and Machine Learning concepts**

- TF 2.0 vs. 1.X comparison
- Name scopes
- Basic Deep Learning Models, k-Means
- Linear and Logistic Regression
- Softmax classification

HW2 (10pts)

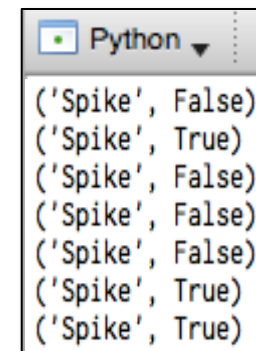
- **Neural Networks 1/2**

- Multi-layer Neural Network
- Gradient descent and Backpropagation

# TensorFlow

- Example:

```
1  ## Detecting Spikes:
2  import tensorflow as tf
3  sess = tf.InteractiveSession()
4
5  # Create a boolean variable called `spike` to detect a sudden increase in
6  # a series of numbers. Since all variables must be initialized, initialize the
7  # variable by calling `run()` on its `initializer`:
8  vector = [3., -2., 8., -4., 0.2, 2.3, 7.5, 14.8]
9  spike = tf.Variable(False)
10
11 # Initiallizing our variable in one of two ways (choose one):
12 # sess.run(spike.initializer)
13 spike.initializer.run()
14
15 # Loop through the data and update the spike variable when there is a
16 # significant increase:
17 for i in range(1, len(vector)):
18     if vector[i] - vector[i-1] > 5:
19         updater = tf.assign(spike, tf.constant(True))
20         updater.eval()
21     else:
22         tf.assign(spike, False).eval()
23     print("Spike", spike.eval())
24
25 # Check to see if there some uninitialized variables:
26 print(sess.run(tf.report_uninitialized_variables()))
27
28 sess.close()
```



```
Python ▼
('Spike', False)
('Spike', True)
('Spike', False)
('Spike', False)
('Spike', False)
('Spike', True)
('Spike', True)
```

Shows the  
difference  
between

... try it in class

# TensorFlow

- Example:

```
1  ## Saving Variables in TensorFlow
2  import tensorflow as tf
3  sess = tf.InteractiveSession()
4
5  # Create a boolean vector called `spike` to locate a sudden spike in data.
6  # Since all variables must be initialized, initialize the variable by calling
7  # `run()` on its `initializer`.
8  vector = [3., -2., 8., -4., 0.2, 2.3, 7.5, 14.8]
9  spikes = tf.Variable([False] * len(vector), name='spikes')
10
11 # Initiallizing our variable in one of two ways (choose one):
12 # sess.run(spikes.initializer)
13 spikes.initializer.run()
14
15 # The saver op will enable saving and restoring
16 saver = tf.train.Saver()
17
18 # Loop through the data and update the spike variable when there is a significant
19 # increase
20 for i in range(1, len(vector)):
21     if vector[i] - vector[i-1] > 5:
22         spikes_val = spikes.eval()
23         spikes_val[i] = True
24         updater = tf.assign(spikes, spikes_val)
25         updater.eval()
26
27 save_path = saver.save(sess, "./spikes.ckpt")
28 print("spikes data saved in file: %s" % save_path)
29 sess.close()
```

spikes data saved in file: ./spikes.ckpt

- checkpoint
- spikes.ckpt.data-00000-of-00001
- spikes.ckpt.index
- spikes.ckpt.meta

... try it in class

# TensorFlow

- Example:

```
1  ## Loading Variables in TensorFlow
2  import tensorflow as tf
3  sess = tf.InteractiveSession()
4
5  # Create a boolean vector called `spike` to locate a sudden spike in data.
6  # Since all variables must be initialized, initialize the variable by calling
7  # `run()` on its `initializer`.
8  spikes = tf.Variable([False]*8, name='spikes')
9  saver = tf.train.Saver()
10
11  saver.restore(sess, "./spikes.ckpt")
12  print(spikes.eval())
13
14  sess.close()
```

```
INFO:tensorflow:Restoring parameters from ./spikes.ckpt
[False False  True False False False  True  True]
```

# TensorFlow

- Example:

```
1  ## Log example:
2  import tensorflow as tf
3
4  # Let's create a simple matrix:
5  matrix = tf.constant([[3., 4.]])
6
7  # Now negate it:
8  negMatrix = tf.negative(matrix)
9
10 # Let's see where each operation is mapped to:
11 with tf.Session(config=tf.ConfigProto(log_device_placement=True)) as sess:
12     result = sess.run(negMatrix)
13
14 # Print results on screen:
15 print(result)
16 print(matrix.shape)
17
18 # To access each member inside a tensor do:
19 print(matrix.shape[0])
```

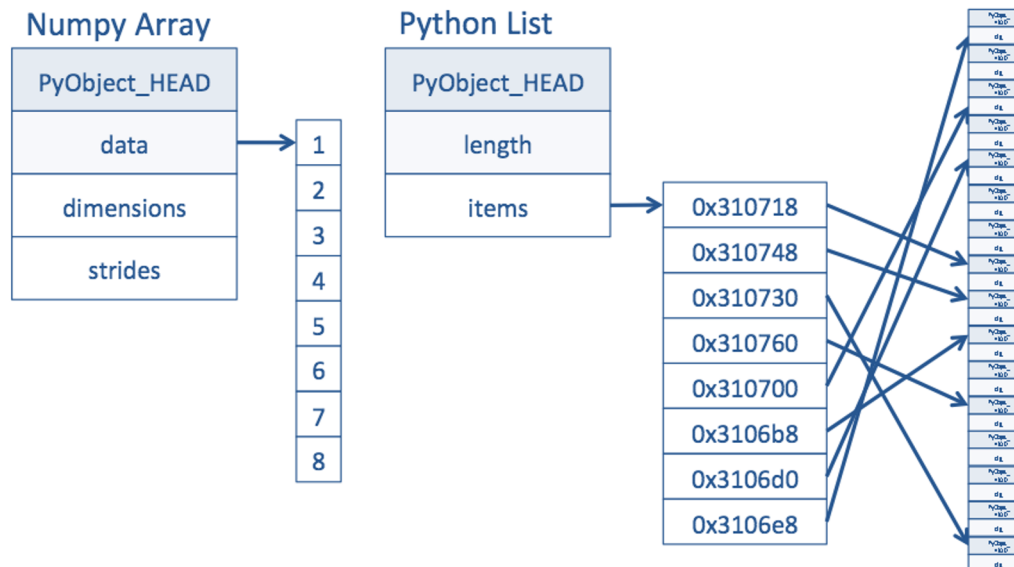
```
2018-05-28 14:51:29.346491: I tensorflow/core/common_runtime/direct_session.cc:297] Device mapping:
2018-05-28 14:51:29.347240: I tensorflow/core/common_runtime/placer.cc:874] Neg: (Neg)/job:localhost/
replica:0/task:0/device:CPU:0
[[-3. -4.]]2018-05-28 14:51:29.347253: I tensorflow/core/common_runtime/placer.cc:874] Const: (Const)
/job:localhost/replica:0/task:0/device:CPU:0
```

```
(1, 2)
1
```

... try it in class

# NumPy arrays recap

- Difference between NumPy arrays vs Python Lists
  - NumPy array:
    - A **NumPy array** is a Python object build around a C array
    - This means that it has a pointer to a **contiguous data buffer of values**
  - Python Lists:
    - A **Python list** has a pointer to a **contiguous buffer of pointers**
    - **All of them point to different Python objects**, which in turn has references to its data (in this case, integers)
  - Conclusion:
    - NumPy is much more efficient than Python, in the **cost of storage** and in **speed of access**



# TensorFlow

*let's recall*


- Let's discuss the basics of **computation graphs** without the context of TensorFlow
- This includes:
  - defining **nodes**
  - defining **edges**
  - **dependencies**
  - examples to illustrate **key principles**



# TensorFlow

*let's recall*

- Graph basics:
  - At **the core** of every TensorFlow program is the **computation graph**
  - It is a specific type of directed graph that is used for **defining computational structure**
  - In TensorFlow it is, a **series of functions chained together**, each passing its output to zero, one, or more functions further along in the chain




$$f(1, 2) = 1 + 2 = 3$$

# TensorFlow

*let's recall*

- Graph basics:
  - **Nodes**: typically drawn as circles, ovals, or boxes, represent some sort of computation or action being done on or with data in the graph's context. In the example below, the operation “add” is the sole **node**.




$$f(1, 2) = 1 + 2 = 3$$

# TensorFlow


*let's recall*

- Graph basics:
  - **Edges**: are the actual **values** that get passed to and from **Operations**, and are typically drawn as **arrows**
  - In the “add” example, the inputs 1 and 2 are both edges leading into the node, while the output 3 is an edge leading out of the node
  - We can think of edges as the link between different **Operations as they carry information** from one node to the next



$$f(1, 2) = 1 + 2 = 3$$

# TensorFlow



# TensorFlow

We can decompose this graphical representation as a series of equations like this:

$$a = \text{input}_1; b = \text{input}_2$$

$$c = a \cdot b; d = a + b$$

$$e = c + d$$


to solve  $e$  for  $a = 5$  and  $b = 3$ ,

$$a = 5; b = 3$$


$$e = (a \cdot b) + (a + b)$$

$$e = (5 \cdot 3) + (5 + 3)$$

$$e = 15 + 8 = 23$$




# TensorFlow



# TensorFlow


- Graph basics:
  - Dependencies:** there are certain types of connections between nodes that aren't allowed, the most common of which is one that creates an unresolved *circular dependency*



# TensorFlow

- Graph basics:
  - Dependencies: ...

let's look at what happens if the multiplication node c is unable to finish its computation (for whatever reason):






# TensorFlow

- Graph basics:
  - Dependencies: ...


What happens if one of the inputs fails to pass its data on to the next functions in the graph?



# TensorFlow

- Graph basics:
  - Dependencies: ...


Let's see what happens if we redirect the output of a graph back into an earlier portion of it:



# TensorFlow

- Graph basics:
  - Because of this, truly circular dependencies can't be expressed in TensorFlow, which is not a bad thing.
- Dependencies: ...

Let's provide an initial state to the value feeding into either *b* or *e*. Let's give the graph a kick-start by giving the output of *e* an initial value of 1:



# TensorFlow

- Graph basics:
  - Dependencies: ...


By unrolling our graph like this, we can simulate useful cyclical dependencies while maintaining a deterministic computation.



# TensorFlow

*let's recall*

- Building our first graph in TensorFlow:



```
simple_graph.py
1 ## Building our first TensorFlow graph:
2
3 # First we need to import TensorFlow:
4 import tensorflow as tf
5
6 # Let's define our input nodes:
7 a = tf.constant(5, name="input_a")
8 b = tf.constant(3, name="input_b")
9
10 # Defining the next two nodes in our graph:
11 c = tf.multiply(a,b, name="mul_c")
12 d = tf.add(a,b, name="add_d")
13
14 # This last line defines the final node in our graph:
15 e = tf.add(c,d, name="add_e")
```

# TensorFlow

*let's recall*

- Building our first graph in TensorFlow:

```
Python /Users/alex
In [6]: whos
Variable Type Data/Info
-----
a Tensor Tensor("input_a_1:0", shape=(), dtype=int32)
b Tensor Tensor("input_b_1:0", shape=(), dtype=int32)
c Tensor Tensor("mul_c:0", shape=(), dtype=int32)
d Tensor Tensor("add_d:0", shape=(), dtype=int32)
e Tensor Tensor("add_e:0", shape=(), dtype=int32)
tf module <module 'tensorflow' from<...>tensorflow/__init__.pyc'>
```

```
simple_graph.py
1 ## Building our first TensorFlow graph:
2
3 # First we need to import TensorFlow:
4 import tensorflow as tf
5
6 # Let's define our input nodes:
7 a = tf.constant(5, name="input_a")
8 b = tf.constant(3, name="input_b")
9
10 # Defining the next two nodes in our graph:
11 c = tf.multiply(a,b, name="mul_c")
12 d = tf.add(a,b, name="add_d")
13
14 # This last line defines the final node in our graph:
15 e = tf.add(c,d, name="add_e")
```

To run we must add the two extra lines and run them in the shell:

```
In [7]: sess = tf.Session()
In [8]: sess.run(e)
Out[8]: 23
```

# TensorFlow

- Let's construct the actual graph using TensorBoard:
  - First, we need to make sure we have generated summary data in a log directory by creating a summary writer:

`tf.Graph().as_default()`

`sess.graph.as_graph_def()`  
contains the graph definition that  
enables the Graph Visualizer

```
21 # To create the graph:  
22 sess.graph.as_graph_def()  
23 file_writer = tf.summary.FileWriter('./', sess.graph)
```

*More on TF 2.x later!*

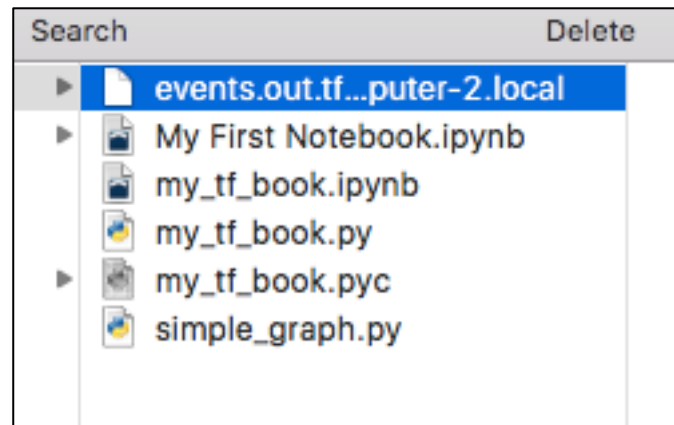
Docstring:  
Writes `Summary` protocol  
buffers to event files.

`'/path/to/logs'`

`tf.summary.FileWriter` is deprecated. Please use `tf.compat.v1.summary.FileWriter` instead.

# TensorFlow

- Let's construct the actual graph using TensorBoard:
- Once we run the previous code, a file with the session is generated in our current folder:

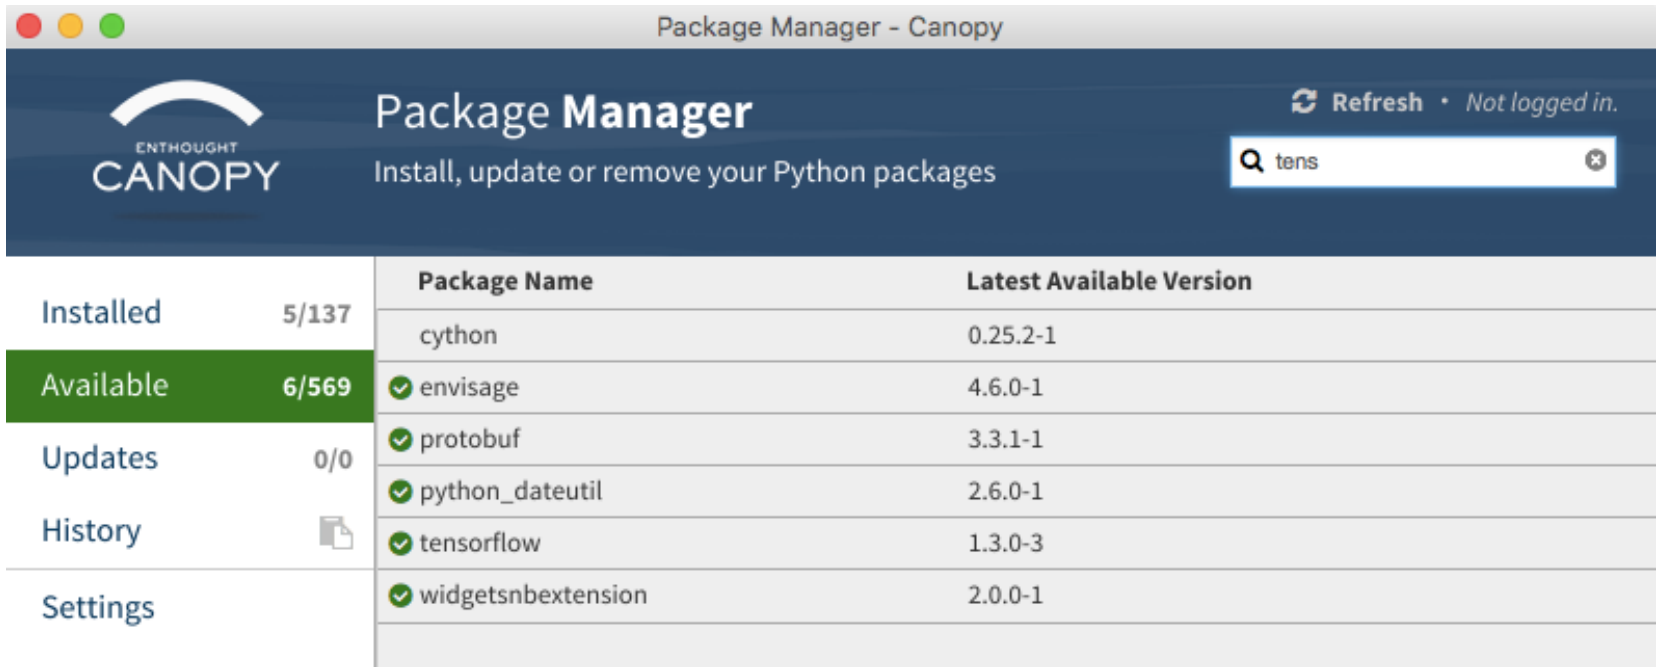




# TensorFlow

*legacy*

- Let's construct the actual graph using TensorBoard:
  - Before we continue, we need to check if we have TensorBoard installed in our system:



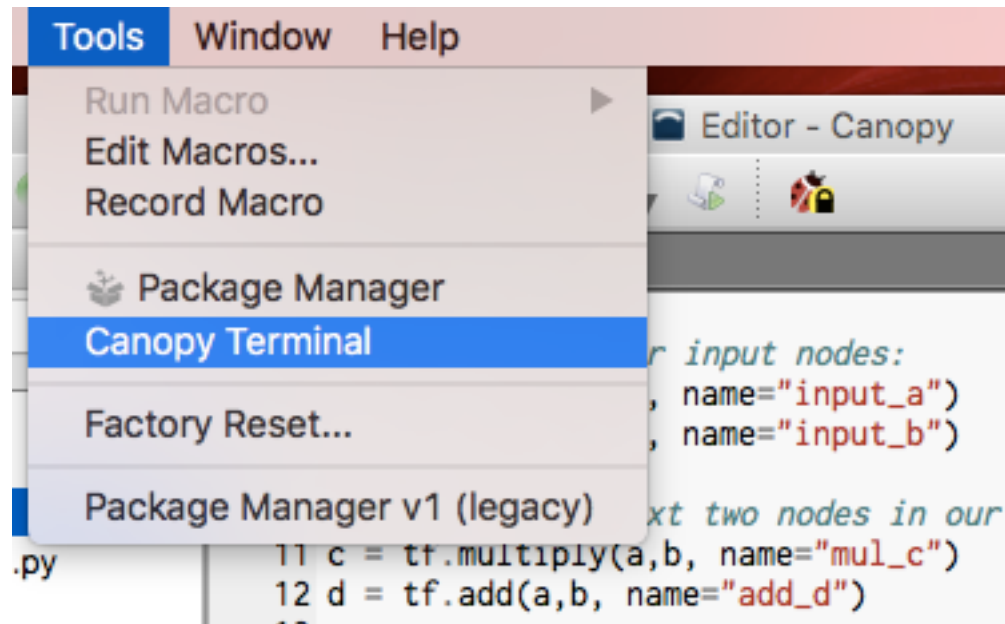
The screenshot shows the Canopy Package Manager window. The title bar reads "Package Manager - Canopy". The interface has a dark blue header with the Canopy logo on the left, the text "Package Manager" and "Install, update or remove your Python packages" in the center, and a "Refresh" button and "Not logged in." status on the right. A search bar contains the text "tens". On the left side, there is a sidebar with links: "Installed 5/137", "Available 6/569" (highlighted in green), "Updates 0/0", "History" (with a document icon), and "Settings". The main area displays a table of installed and available packages.

Package Name	Latest Available Version
cython	0.25.2-1
✓ envisage	4.6.0-1
✓ protobuf	3.3.1-1
✓ python_dateutil	2.6.0-1
✓ tensorflow	1.3.0-3
✓ widgetsnbextension	2.0.0-1

# TensorFlow

*legacy*

- Let's construct the actual graph using TensorBoard:
- Canopy **does not provide TensorBoard** in its repository, therefore we need to install it via the Canopy Terminal:



# TensorFlow

*legacy*

- Let's construct the actual graph using TensorBoard:
- WARNING:** We check for compatibility: **tensorflow** + **tensorboard**

```
tensorboard 1.8.0
tensorflow 1.6.0
```

When installing tensorboard you might get this message:

```
tensorflow 1.6.0 has requirement tensorboard<1.7.0,>=1.6.0, but you'll have tensorboard 1.8.0 which is incompatible.
```

When running tensorboard you might get this message:

```
(Canopy 64bit) Alexandomputer2:alex_examples alex$ tensorboard --logdir ./ --event_file events.out.tfevents.1527382420.Alexandomputer2
2018-05-26 18:00:37.024647: I tensorflow/core/platform/cpu_feature_guard.cc:140]
Your CPU supports instructions that this TensorFlow binary was not compiled to
use: SSE4.1
```

So you have to install a compatible version:

```
alex$ pip install tensorboard==1.6.0
```

Make sure they match:

```
tensorboard 1.6.0
tensorflow 1.6.0
```

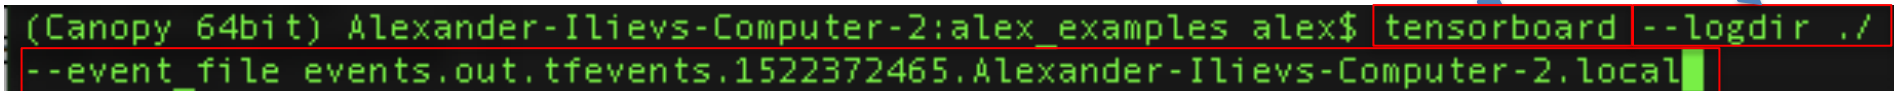
When running tensorboard it runs without warnings:

```
(Canopy 64bit) Alexandomputer2:alex_examples alex$ tensorboard --logdir ./ --event_file events.out.tfevents.1527382420.Alexandomputer2
TensorBoard 1.6.0 at http://Alexandomputer2:6006 (Press CTRL+C to quit)
```

# TensorFlow

*legacy*

- Let's construct the actual graph using TensorBoard:
  - Once we have the event file(s), we run TensorBoard while providing the log directory:

A terminal window showing a command to run TensorBoard. The command is: `(Canopy 64bit) Alexander-Ilievs-Computer-2:alex_examples alex$ tensorboard --logdir ./ --event_file events.out.tfevents.1522372465.Alexander-Ilievs-Computer-2.local`. Three blue arrows point from the text above to parts of the command: one from 'TensorBoard' to 'tensorboard', one from 'log directory' to '--logdir ./', and one from 'event file(s)' to '--event\_file events.out.tfevents.1522372465.Alexander-Ilievs-Computer-2.local'.

```
(Canopy 64bit) Alexander-Ilievs-Computer-2:alex_examples alex$ tensorboard --logdir ./ --event_file events.out.tfevents.1522372465.Alexander-Ilievs-Computer-2.local
```

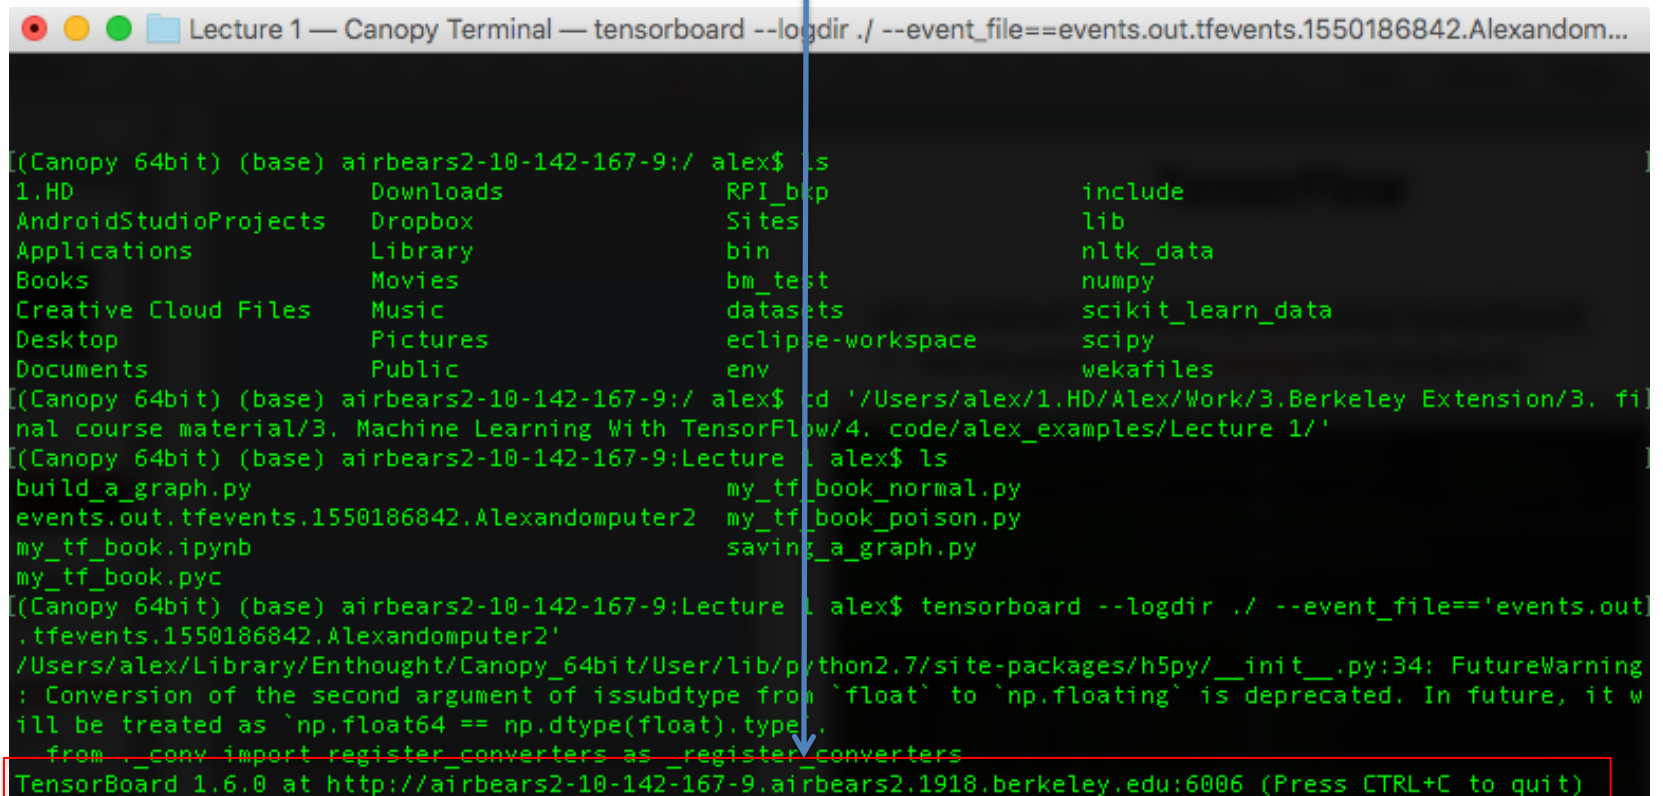
- and specifically request the file to be executed
- or if you are in the same directory where the file resides simply run:

```
tensorboard --logdir ./
```

# TensorFlow

*legacy*

- Let's construct the actual graph using TensorBoard:
  - And the graph is already **running** in the background:



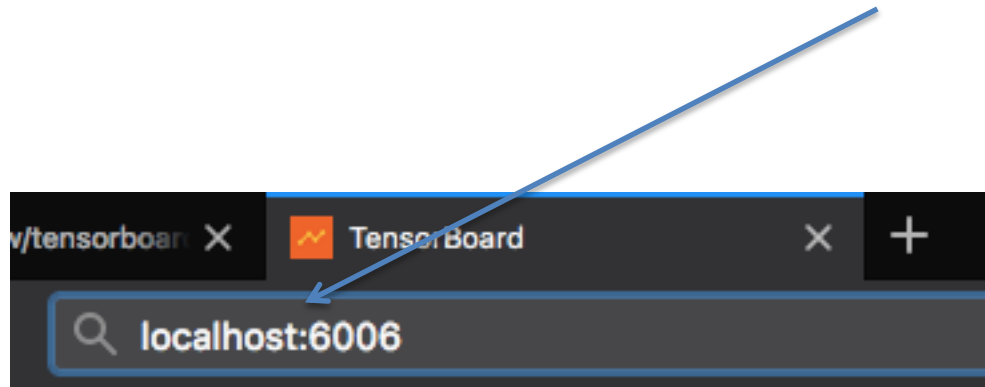
The screenshot shows a Canopy Terminal window titled "Lecture 1 — Canopy Terminal — tensorboard --logdir ./ --event\_file==events.out.tfevents.1550186842.Alexandom...". The terminal output shows the user navigating to the directory where the graph was saved and running the `tensorboard` command. The output of the command is highlighted with a red box:

```
TensorBoard 1.6.0 at http://airbears2-10-142-167-9.airbears2.1918.berkeley.edu:6006 (Press CTRL+C to quit)
```

# TensorFlow

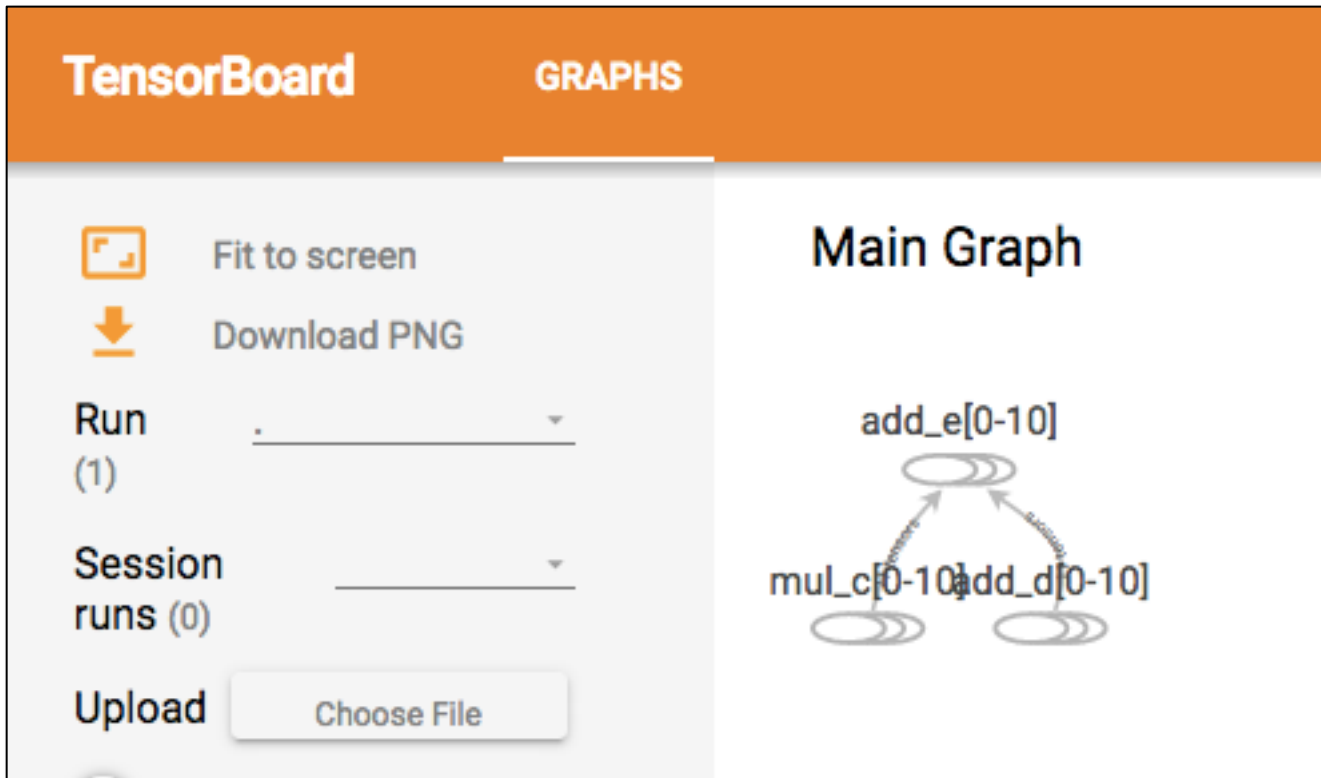
*legacy*

- Let's construct the actual graph using TensorBoard:
- Next, we open our browser and **look at the graph by typing:**



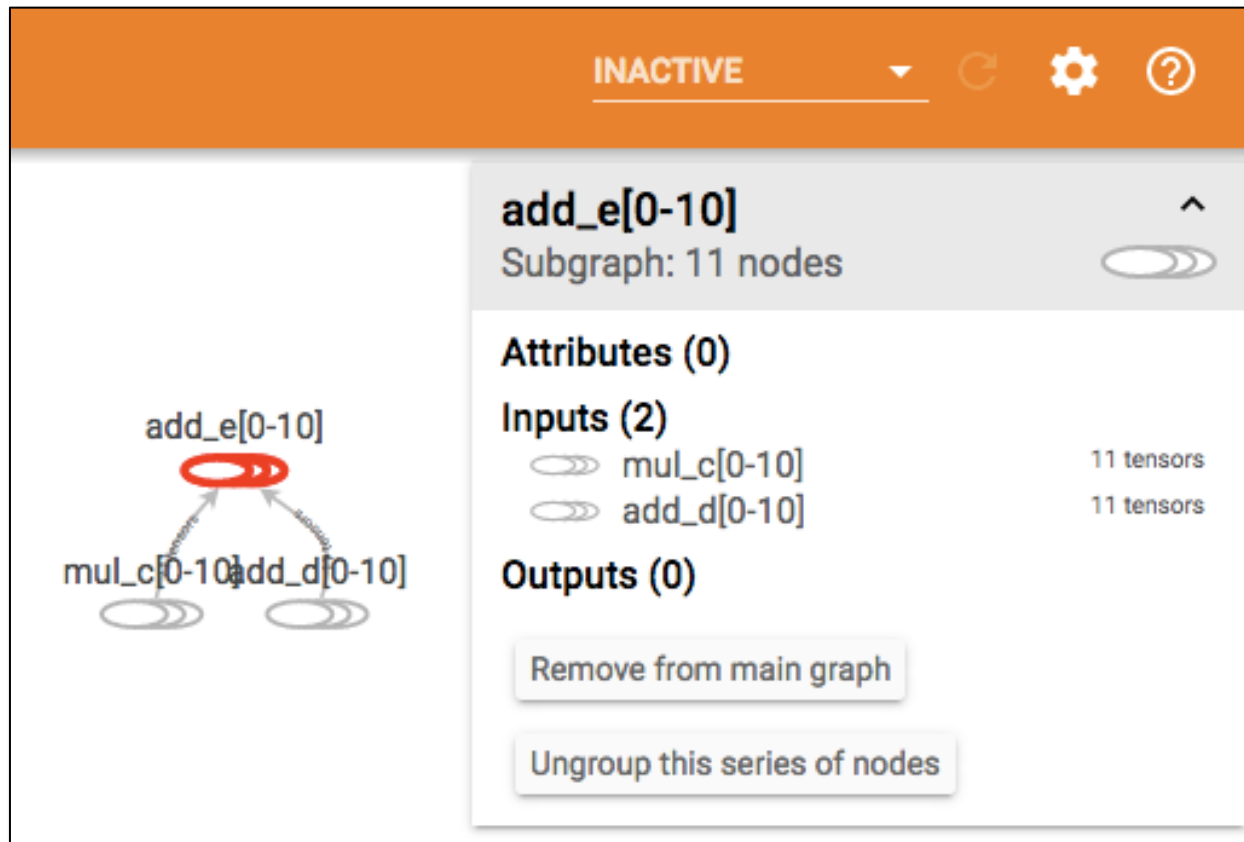
# TensorFlow

- Let's construct the actual graph using TensorBoard:
- Below is a graphical representation of our first **graph**:



# TensorFlow

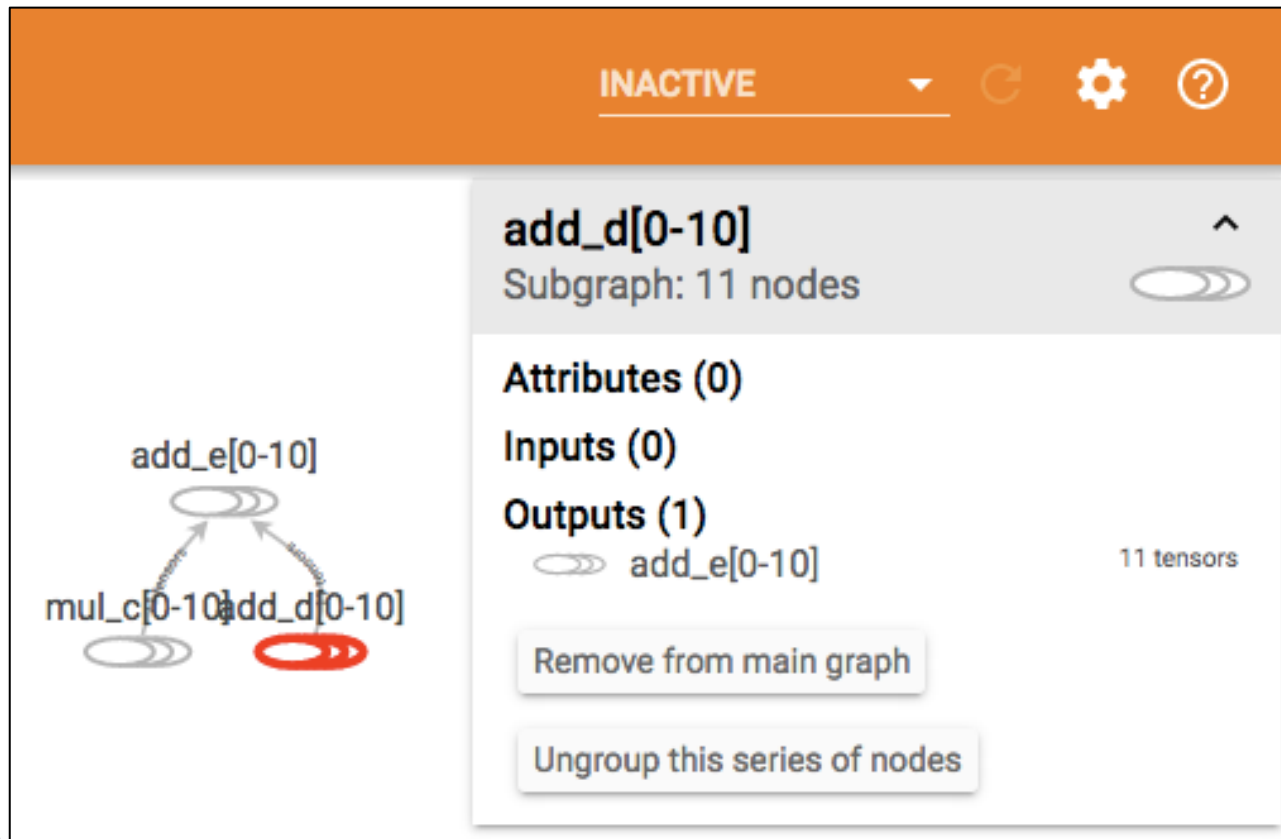
- Let's construct the actual graph using TensorBoard:
- Simply click on any of the **nodes** to inspect them more closely:





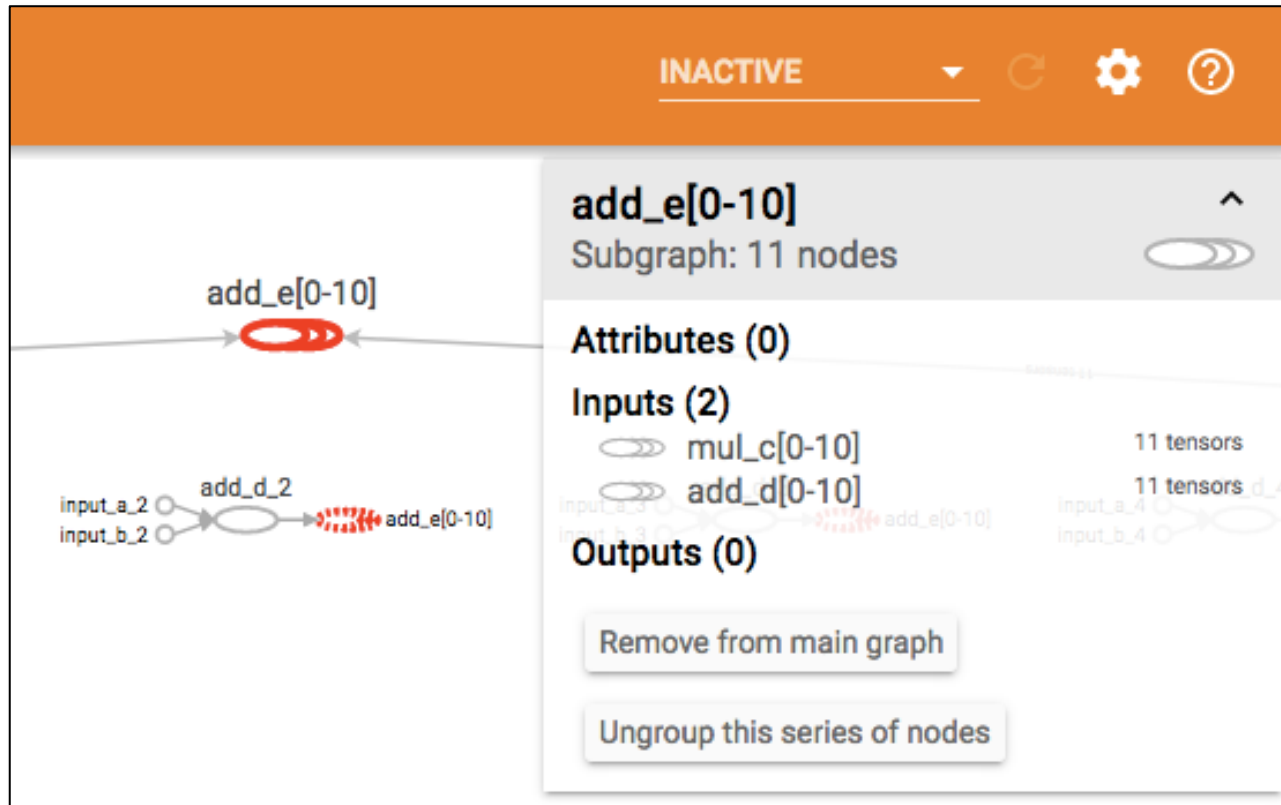
# TensorFlow

- Let's construct the actual graph using TensorBoard:
- Simply click on any of the **nodes** to inspect them more closely:












# TensorFlow

- Let's construct the actual graph using TensorBoard:
- Simply click on any of the **nodes** to inspect them more closely:



# TensorFlow


Symbol	Meaning
	<i>High-level</i> node representing a name scope. Double-click to expand a high-level node.
	Sequence of numbered nodes that are not connected to each other.
	Sequence of numbered nodes that are connected to each other.
	An individual operation node.
	A constant.
	A summary node.
	Edge showing the data flow between operations.
	Edge showing the control dependency between operations.
	A reference edge showing that the outgoing operation node can mutate the incoming tensor.

# TensorFlow

*legacy*

- Let's construct the actual graph using TensorBoard:
  - Once we are done constructing our graph, we need to clean up and close the *file\_writer* and *sess*:

```
25 # We clean up before we exit:  
26 file_writer.close()  
27 sess.close()
```

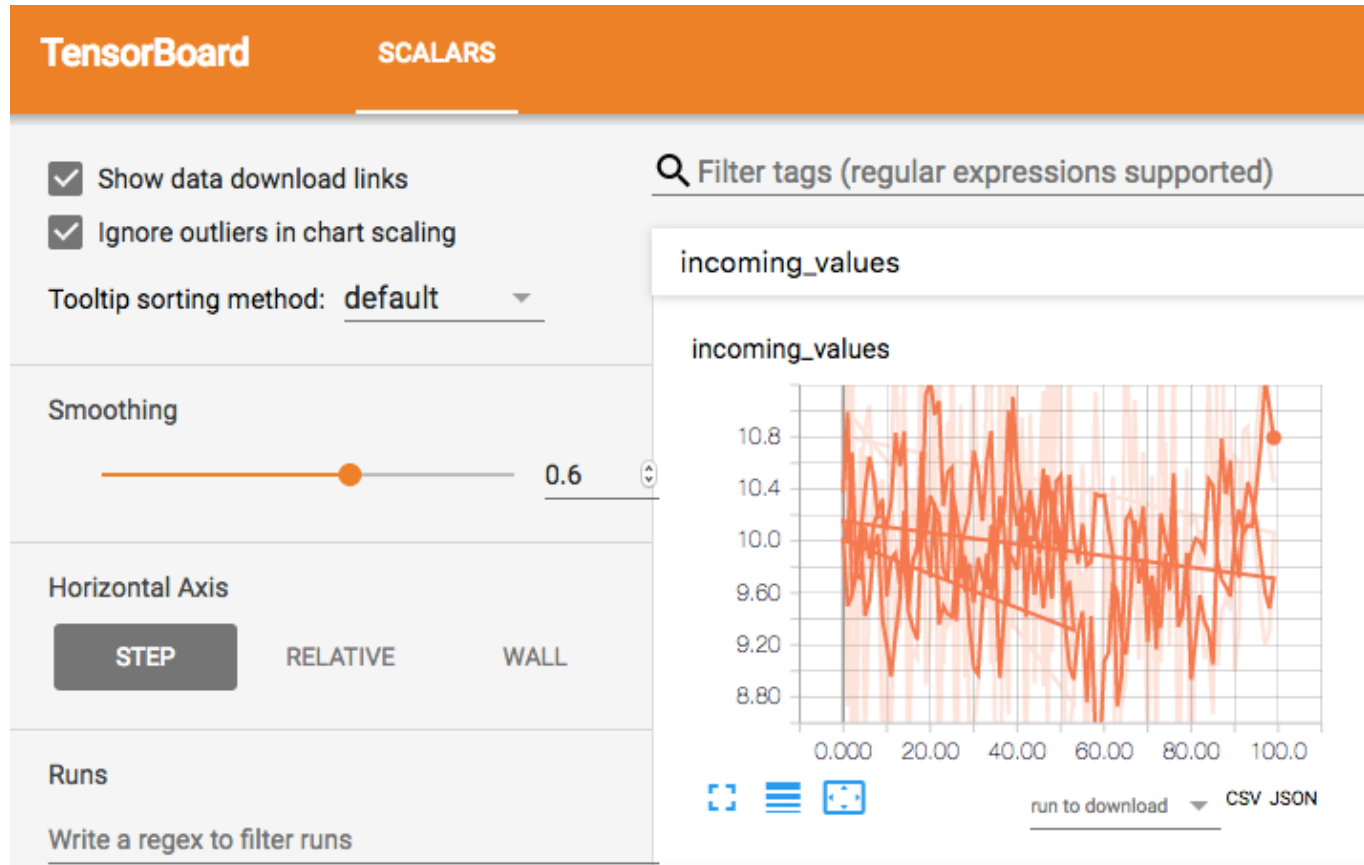


*tf.summary.FileWriter is deprecated. Please use tf.compat.v1.summary.FileWriter instead.*

- In general, *Session* objects close automatically when the program terminates (or, in the interactive case, when you close/restart the Python kernel)
- However, **it's best to explicitly close out** of the *Session* to avoid any sort of weird edge case scenarios.

# TensorFlow

- Let's construct the actual graph using TensorBoard:



	Name	Smoothed	Value	Step	Time	Relative
	incoming_aver	10.06	9.825	67.00	Sun Jun 3, 18:38:39	8m 58s

# TensorFlow operations

*recall*

- Remember: **Tensors** are just a superset of **matrices**!
- **TensorFlow** Operations, aka Ops, **are nodes** that perform computations on or with Tensor objects
- After computation, they return zero or more tensors, which can be used by other Ops later in the graph
- **To create an Operation**, you **call its constructor** in Python
- The Python **constructor returns a handle** to the Operation's output, then it is passed on to other **Ops** or **Session.run**

# TensorFlow operations

*recall*

- Example:

```
import tensorflow as tf
import numpy as np

# Initialize some tensors to use in computation
a = np.array([2, 3], dtype=np.int32)
b = np.array([4, 5], dtype=np.int32)

# Use `tf.add()` to initialize an "add" Operation
# The variable `c` will be a handle to the Tensor output of
this Op
c = tf.add(a, b)
```

# TensorFlow operations

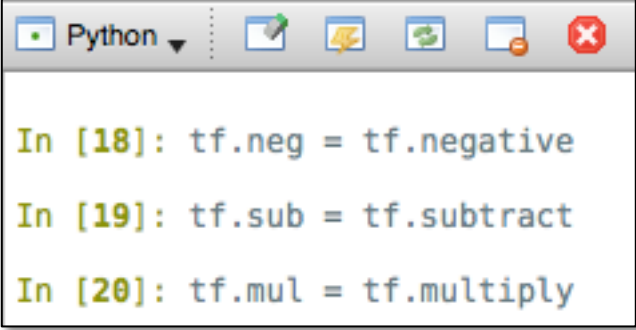
- Overloaded operators:
  - TensorFlow also **overloads** common mathematical operators to make **multiplication**, **addition**, **subtraction**, and **other operations** more concise
  - If one or more arguments to the operator is a Tensor object, a **TensorFlow Operation** will be called and **added to the graph**
  - For example, you can easily add two tensors together like this:

```
# Assume that `a` and `b` are `Tensor` objects with  
matching shapes  
c = a + b
```



# TensorFlow operations

- Creating aliases:
  - TensorFlow allows us to create **aliases** for all common mathematical operators such as **negation**, **subtraction**, **multiplication**, or **other operations** to be more concise
  - For example, you can easily create aliases like this:



```
In [18]: tf.neg = tf.negative
In [19]: tf.sub = tf.subtract
In [20]: tf.mul = tf.multiply
```

# TensorFlow operations

- Overloaded operators using aliases (from previous slide):
  - A complete list of overloaded operators is given for tensors:

## Unary operators

Operator	Related Tensor-Flow Operation	Description
<code>-x</code>	<code>tf.neg()</code> <code>tf.neg = tf.negative</code>	Returns the negative value of each element in <code>x</code>
<code>~x</code>	<code>tf.logical_not()</code>	Returns the logical NOT of each element in <code>x</code> . Only compatible with <code>Tensor</code> objects with <code>dtype</code> of <code>tf.bool</code>
<code>abs(x)</code>	<code>tf.abs()</code>	Returns the absolute value of each element in <code>x</code>

# TensorFlow operations

- Overloaded operators using aliases:

Binary operators	Related TensorFlow Operation	Description
$x + y$	<code>tf.add()</code>	Add x and y, element-wise
$x - y$	<code>tf.sub()</code> <code>tf.sub = tf.subtract</code>	Subtract y from x, element-wise
$x * y$	<code>tf.mul()</code> <code>tf.mul = tf.multiply</code>	Multiply x and y, element-wise
$x / y$ (Python 2)	<code>tf.div()</code>	Will perform element-wise integer division when given an integer type tensor, and floating point (“true”) division on floating point tensors
$x / y$ (Python 3)	<code>tf.truediv()</code>	Element-wise floating point division (including on integers)
$x // y$ (Python 3)	<code>tf.floordiv()</code>	Element-wise floor division, not returning any remainder from the computation

# TensorFlow operations

- Overloaded operators:

Binary operators	Related TensorFlow Operation	Description
<code>x % y</code>	<code>tf.mod()</code>	Element-wise modulo
<code>x ** y</code>	<code>tf.pow()</code>	The result of raising each element in <code>x</code> to its corresponding element <code>y</code> , element-wise
<code>x &lt; y</code>	<code>tf.less()</code>	Returns the truth table of <code>x &lt; y</code> , element-wise
<code>x &lt;= y</code>	<code>tf.less_equal()</code>	Returns the truth table of <code>x &lt;= y</code> , element-wise
<code>x &gt; y</code>	<code>tf.greater()</code>	Returns the truth table of <code>x &gt; y</code> , element-wise
<code>x &gt;= y</code>	<code>tf.greater_equal()</code>	Returns the truth table of <code>x &gt;= y</code> , element-wise

# TensorFlow operations

- Overloaded operators:
  - A complete list of overloaded operators is given for tensors:

Binary operators	Related TensorFlow Operation	Description
$x \& y$	<code>tf.logical_and()</code>	Returns the truth table of $x \& y$ , element-wise. <code>dtype</code> must be <code>tf.bool</code>
$x   y$	<code>tf.logical_or()</code>	Returns the truth table of $x   y$ , element-wise. <code>dtype</code> must be <code>tf.bool</code>
$x \wedge y$	<code>tf.logical_xor()</code>	Returns the truth table of $x \wedge y$ , element-wise. <code>dtype</code> must be <code>tf.bool</code>

A quick reminder:

$A$	$B$	AND	OR	XOR	NOT $B$
0	0	0	0	0	1
0	1	0	1	1	0
1	0	0	1	1	1
1	1	1	1	0	0

# TensorFlow operations

- Overloaded operators:
  - Using these overloaded operators **can be great** when quickly putting together code
  - Technically, the **==** operator is overloaded as well, **but it will not return a Tensor of boolean values**. It will return **True** if the two tensors being compared are the same object, and **False** otherwise.
  - To check for **equality** or **inequality**, try:  
**tf.equal()** and **tf.not\_equal**, respectively.

# TensorFlow graphs

- Creating a Graph is simple:
  - its constructor doesn't take any variables:

```
import tensorflow as tf

# Create a new graph:
g = tf.Graph()
```

- Once we have our Graph initialized, we can add Operations to it:

```
with g.as_default():
    # Create Operations as usual; they will be added to graph
    `g`
    a = tf.mul(2, 3)
    ...
```

# TensorFlow graphs

- TensorFlow **automatically creates a Graph** when the library is loaded and assigns it to be the **default**.
- Thus, any Operations, tensors, etc. defined outside of a **Graph.as\_default()** context manager will automatically be placed in the default graph:

```
# Placed in the default graph
in_default_graph = tf.add(1,2)

# Placed in graph `g`
with g.as_default():
    in_graph_g = tf.mul(2,3)

# We are no longer in the `with` block, so this is placed in
the default graph
also_in_default_graph = tf.sub(5,1)
```




# TensorFlow graphs

- If you'd like to get a handle to the default graph, use the `tf.get_default_graph()` function:

```
default_graph = tf.get_default_graph()
```

- In most TensorFlow programs, you will only ever deal with the default graph
- When defining multiple graphs in one file, it's better to either not use the default graph or immediately assign a handle to it
- This ensures that nodes are added to each graph in a uniform manner



```
tf.compat.v1.get_default_graph()
```

# TensorFlow graphs

- This ensures that nodes are added to each graph in a uniform manner:

**Correct - Create new graphs, ignore default graph:**

```
import tensorflow as tf

g1 = tf.Graph()
g2 = tf.Graph()

with g1.as_default():
    # Define g1 Operations, tensors, etc.
    ...

with g2.as_default():
    # Define g2 Operations, tensors, etc.
    ...
```

# TensorFlow graphs

- This ensures that nodes are added to each graph in a uniform manner:

**Correct - Get handle to default graph:**

```
import tensorflow as tf

g1 = tf.get_default_graph()
g2 = tf.Graph()
    tf.compat.v1.get_default_graph()

with g1.as_default():
    # Define g1 Operations, tensors, etc.
    ...

with g2.as_default():
    # Define g2 Operations, tensors, etc.
    ...
```

# TensorFlow graphs

- This ensures that nodes are added to each graph in a uniform manner:

**Incorrect - Mix default graph and user-created graph styles:**

```
import tensorflow as tf

g2 = tf.Graph()

# Define default graph Operations, tensors, etc.
...

with g2.as_default():
    # Define g2 Operations, tensors, etc.
    ...
```

# TensorFlow graphs

- Additionally, it is possible to load in previously defined models from other **TensorFlow** scripts and assign them to **Graph** objects
- This can be done by using a combination of the **graph.as\_graph\_def()** and **tf.import\_graph\_def** functions
- Thus, a user can compute and use the output of several separate models in the same Python file.

# TensorFlow sessions

- **Sessions**, are responsible for **graph execution**
- The **constructor** takes in three optional parameters:

**target** specifies the execution engine to use:

- When using sessions in a distributed setting, this parameter is used to connect to `tf.train.Server` instances

**graph** specifies the **Graph** object that will be launched in the **Session**.

- When using multiple graphs, it's best to explicitly pass in the **Graph** you'd like to run (instead of creating the **Session** inside of a **with** block).

**config** allows users to specify options to configure the session, such as limiting the **number of CPUs or GPUs to use**, setting optimization parameters for graphs, and logging options.

# TensorFlow sessions

- In a typical **TensorFlow** program, **Session** objects will be created without changing any of the default construction parameters:

```
import tensorflow as tf

# Create Operations, Tensors, etc (using the default graph)
a = tf.add(2, 5)
b = tf.mul(a, 3)

# Start up a `Session` using the default graph
sess = tf.Session()
```

- Note that these two calls are identical:

```
sess = tf.Session()
sess = tf.Session(graph=tf.get_default_graph())
```

The name `tf.Session` is deprecated. Please use `tf.compat.v1.Session` instead.

# TensorFlow sessions

- Once a **Session** is opened, you can use its primary method, **run()**, to calculate the value of a desired Tensor output:

```
sess.run(b)  # Returns 21
```

- **Session.run()** takes in one required parameter, **fetches**,  
(as well as **three optional** parameters: **feed\_dict**, **options**, and **run\_metadata**)
- In the previous example, we set **fetches** to the tensor **b**  
(the output of the **tf.mul** Operation)



# TensorFlow sessions

- We can also pass in a list of graph elements:

```
sess.run([a, b]) # returns [7, 21]
```

- When **fetches** is a list, the output of **run()** will be a list with values corresponding to the output of the requested elements.
- In this example, we ask for the values of **a** and **b**, in that order
- Since both **a** and **b** are tensors, we receive their values as output

# TensorFlow sessions

- We can give **fetches** a direct handle to an **Operation**
- An example of this is **tf.global\_variables\_initializer()**, which prepares **all TensorFlow Variable objects** to be used
- We still pass the Op as the **fetches** parameter, but the result of **Session.run()** will be **None**:

```
# Performs the computations needed to initialize Variables,  
but returns `None`
```

```
sess.run(tf.initialize_all_variables())
```

→ `tf.global_variables_initializer()`

*notice the change in latest 1.x TF versions*

*notice the change in TF 2.1*

→ `tf.compat.v1.initialize_all_variables()`

# TensorFlow sessions

- The parameter **feed\_dict** is used to override **Tensor** values in the graph, and it expects a Python dictionary object as input.
- The **keys** in the dictionary **are handles to Tensor objects** that should be overridden, while the **values** can be **numbers, strings, lists, or NumPy arrays** (as described previously)
- The **values must be of the same type** (or able to be converted to the same type) as the Tensor key.

# TensorFlow sessions

- Here is an example of how `feed_dict` is used to override `Tensor` value `a`:

```
import tensorflow as tf
```

```
# Create Operations, Tensor
```

```
a = tf.add(2, 5)
```

```
b = tf.mul(a, 3)
```

```
# Start up a `Session` using
```

```
sess = tf.Session()
```

```
tf.compat.v1.Session()
```

```
# Define a dictionary that  
with 15
```

```
replace_dict = {a: 15}
```

```
# Run the session, passing in `replace_dict` as the value to  
`feed_dict`
```

```
sess.run(b, feed_dict=replace_dict) # returns 45
```

result:

```
In [3]: a = tf.add(2, 5)
```

```
In [4]: b = tf.multiply(a, 3)
```

```
In [5]: sess = tf.Session()
```

```
In [6]: replace_dict = {a: 15}
```



```
In [7]: sess.run(b, feed_dict = replace_dict)
```

```
Out[7]: 45
```

# Machine Learning With TensorFlow

## Class exercise 1/2:

- Given what we discussed in class today create a short TF program that:
- Creates and plots a *normally distributed* cluster of random numbers, say  $[2, 20] \times [\text{var}, \#]$
- Creates and plots two clusters using *poisson distribution*, say  $[6, 20], [2, 20]$  in red
- Comment on your results



# Machine Learning With TensorFlow


```
1 # Normal distribution:
2 import matplotlib.pyplot as plt
3 import tensorflow as tf
4
5 # This is only for Jupyter plotting in-line:
6 # % matplotlib inline
7
8 a = tf.random_normal([2,20])
9
10 sess = tf.Session()
11 out = sess.run(a)
12 x, y = out
13 plt.scatter(x, y)
14
15 plt.show()
```

/2:

ate a short TF program that:

f random numbers, say [2, 20] \*[var,#]

bution, say [6,20],[2,20] in red



```
1 # Poisson distribution:
2 import matplotlib.pyplot as plt
3 import tensorflow as tf
4
5 # This is only for Jupyter plotting in-line:
6 # % matplotlib inline
7
8 a = tf.random_poisson([6,20],[2,20])
9
10 sess = tf.Session()
11 out = sess.run(a)
12 x, y = out
13
14 color = 'red'
15 plt.scatter(x, y, c=color, label=color, alpha=0.35, edgecolors='black')
16
17 plt.show()
```

# TensorFlow placeholder

- Adding Inputs with **placeholder** nodes
- To take values from the client and plug them into our graph we use what is called a “**placeholder**”
- **Placeholders**, act as if they are Tensor objects, but they do not have their values specified when created
- Instead, **they hold the place for a Tensor** that will be fed at runtime, hence **become input nodes**

# TensorFlow placeholder

- Adding Inputs with **placeholder** nodes
- Example:

```
import tensorflow as tf
import numpy as np

# Creates a placeholder vector of length 2 with data type int32
a = tf.placeholder(tf.int32, shape=[2], name="my_input")

# Use the placeholder as if it were any other Tensor object
b = tf.reduce_prod(a, name="prod_b")
c = tf.reduce_sum(a, name="sum_c")

# Finish off the graph
d = tf.add(b, c, name="add_d")
```

```
tf.math.reduce_prod()
tf.math.reduce_sum()
tf.math.add()
```



# TensorFlow placeholder

- Adding Inputs with **placeholder** nodes
- **tf.placeholder** takes in a required parameter **dtype**, as well as the optional parameter **shape**:
  - **dtype** specifies the data type of values that will be passed into the placeholder. This is required, in order to ensure that there will be **no type mismatch errors**
  - **shape** specifies what shape the fed Tensor will be. The default value of **shape** is **None**, which means a Tensor of any shape will be accepted

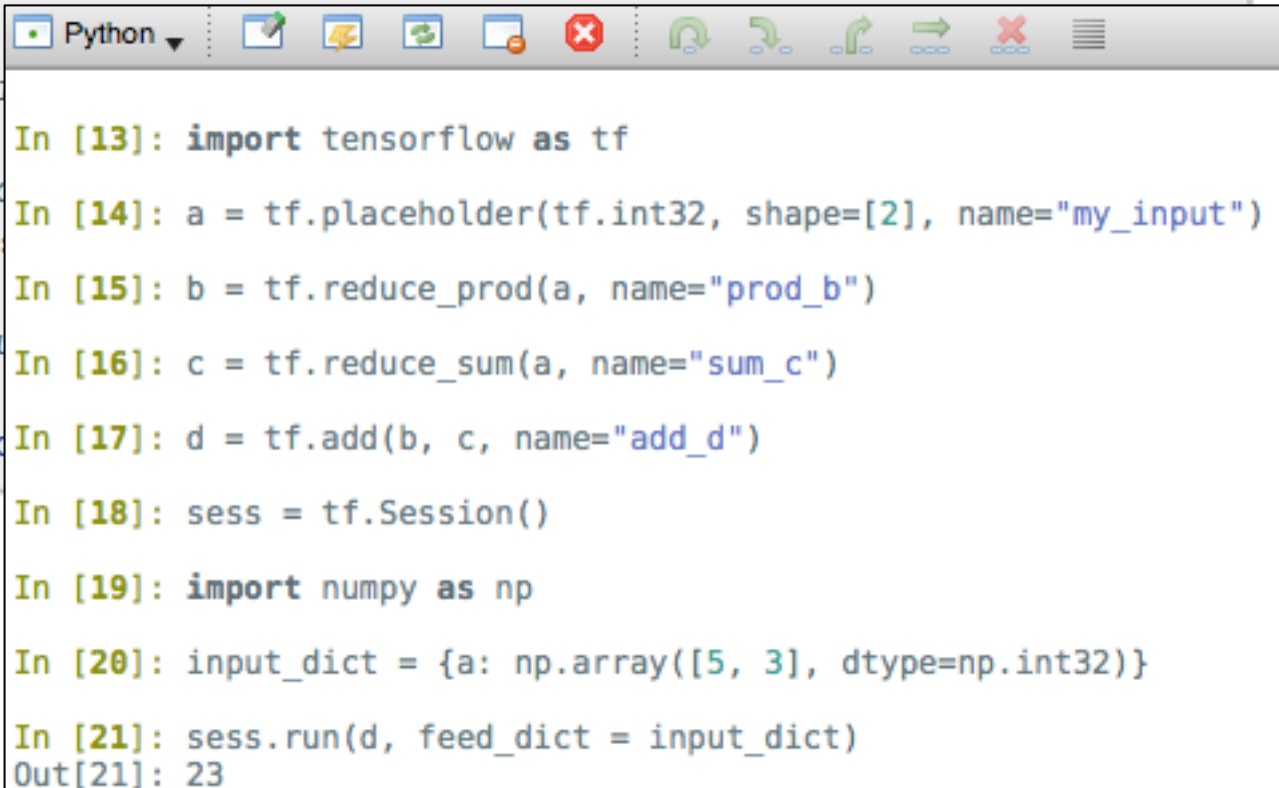
# TensorFlow placeholder

- You can also specify a name identifier to `tf.placeholder`

```
# Open a TensorFlow Session  
sess = tf.Session()
```

result:

```
# Create a dict  
# Key: `a`, the  
# Value: A vector  
input_dict = {a:  
  
# Fetch the value  
into `a`  
sess.run(d, feed
```



The image shows a Jupyter Notebook window with a toolbar at the top containing icons for saving, undo, redo, and other actions. The notebook contains the following code and output:

```
In [13]: import tensorflow as tf  
In [14]: a = tf.placeholder(tf.int32, shape=[2], name="my_input")  
In [15]: b = tf.reduce_prod(a, name="prod_b")  
In [16]: c = tf.reduce_sum(a, name="sum_c")  
In [17]: d = tf.add(b, c, name="add_d")  
In [18]: sess = tf.Session()  
In [19]: import numpy as np  
In [20]: input_dict = {a: np.array([5, 3], dtype=np.int32)}  
In [21]: sess.run(d, feed_dict = input_dict)  
Out[21]: 23
```

# TensorFlow variables

- Tensor and Operation objects are **immutable**
- We need a way to save changing values over time
- This is accomplished in TensorFlow with **Variable objects**
- You can create a **Variable** by using its constructor, `tf.Variable()`:

```
import tensorflow as tf

# Pass in a starting value of three for the variable
my_var = tf.Variable(3, name="my_variable")
```

# TensorFlow variables

- **Variables** can be used in TensorFlow functions/Operations anywhere you might use a **Tensor**
- Its present value will be passed on to the **Operation** using it:

```
add = tf.add(5, my_var)  
mul = tf.mul(8, my_var)
```

- The **initial value** of Variables will often be large tensors of **zeros**, **ones**, or **random values**
- TensorFlow has a number of helper Ops, such as:  
**tf.zeros()**, **tf.ones()**, **tf.random\_normal()**, and **tf.random\_uniform()**

# TensorFlow variables

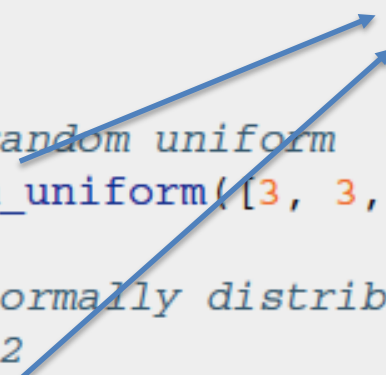
- Each of these takes in a shape parameter which specifies the dimension of the desired Tensor:

```
# 2x2 matrix of zeros  
zeros = tf.zeros([2, 2])
```

```
# vector of length 6 of ones  
ones = tf.ones([6])
```

```
# 3x3x3 Tensor of random uniform values between 0 and 10  
uniform = tf.random_uniform([3, 3, 3], minval=0, maxval=10)
```

```
# 3x3x3 Tensor of normally distributed numbers; mean 0 and  
standard deviation 2  
normal = tf.random_normal([3, 3, 3], mean=0.0, stddev=2.0)
```




```
tf.random.normal()  
tf.random.uniform()
```

# TensorFlow variables

- Instead of using `tf.random_normal()`, you'll often see use of `tf.truncated_normal()` instead

```
# No values below 3.0 or above 7.0 will be returned in this  
Tensor
```

```
trunc = tf.truncated_normal([2, 2], mean=5.0, stddev=1.0)
```



```
tf.random.truncated_normal()
```

- You can pass in these Operations as the initial values of Variables as you would a handwritten Tensor:

```
# Default value of mean=0.0
```

```
# Default value of stddev=1.0
```

```
random_var = tf.Variable(tf.truncated_normal([2, 2]))
```

# TensorFlow variables

- Variable objects **live in the Graph** like most other TensorFlow objects, but their state is actually **managed by a Session**.
- Because of this, **Variables** have an **extra step** involved in order to use them:
  - you **must initialize** the **Variable** within a **Session**

```
init = tf.initialize_all_variables()  
sess = tf.Session()  
sess.run(init)
```

Use `'tf.global_variables_initializer'` instead.

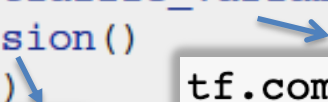
`tf.compat.v1.Session()`

`tf.compat.v1.initialize_all_variables()`

# TensorFlow **changing** variables

- If you'd only like to **initialize a subset** of Variables defined in the graph, you can use **tf.variables\_initializer()**
- This takes in a list of Variables to be initialized:

```
var1 = tf.Variable(0, name="initialize_me")
var2 = tf.Variable(1, name="no_initialization")
init = tf.initialize_variables([var1], name="init_var1")
sess = tf.Session()
sess.run(init)
```



```
tf.compat.v1.Session()
```

```
tf.compat.v1.variables_initializer()
```

- In order to change the value of the **Variable**, you can use the **Variable.assign()** method, which gives the **Variable** the new value

\* Note that **Variable.assign()** is an **Operation**, and must be run in a **Session** to take effect



# TensorFlow **changing** variables

- Example:

```
# Create variable with starting value of 1
my_var = tf.Variable(1)

# Create an operation that multiplies the variable by 2 each
time it is run
my_var_times_two = my_var.assign(my_var * 2)

# Initialization operation
init = tf.initialize_all_variables()

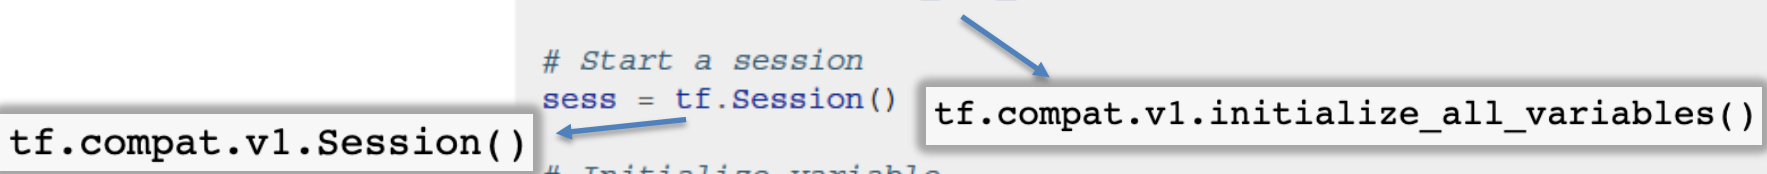
# Start a session
sess = tf.Session()

# Initialize variable
sess.run(init)

# Multiply variable by two and return it
sess.run(my_var_times_two)
## OUT: 2

# Multiply again
sess.run(my_var_times_two)
## OUT: 4

# Multiply again
sess.run(my_var_times_two)
## OUT: 8
```



tf.compat.v1.Session()

tf.compat.v1.initialize\_all\_variables()

# TensorFlow **changing** variables

- For simple incrementing and decrementing of Variables, TensorFlow includes the **Variable.assign\_add()** **Variable.assign\_sub()** methods:

```
# Increment by 1  
sess.run(my_var.assign_add(1))  
  
# Decrement by 1  
sess.run(my_var.assign_sub(1))
```

```
tf.compat.v1.assign_add()  
tf.compat.v1.assign_sub()
```

- Why is this good? ... Because **Sessions** maintain **Variable** values separately, **each Session can have its own current value** for a **Variable** defined in a graph (... on next slide)

# TensorFlow **changing** variables

```
# Create Ops
my_var = tf.Variable(0)
init = tf.initialize_all_variables()

# Start Sessions
sess1 = tf.Session()
sess2 = tf.Session()

# Initialize Variable in sess1, and increment value of my_var
# in that Session
sess1.run(init)
sess1.run(my_var.assign_add(5))
## OUT: 5

# Do the same with sess2, but use a different increment value
sess2.run(init)
sess2.run(my_var.assign_add(2))
## OUT: 2

# Can increment the Variable values in each Session independ-
# ently

sess1.run(my_var.assign_add(5))
## OUT: 10

sess2.run(my_var.assign_add(2))
## OUT: 4
```

Annotations for the code above:

- `tf.initialize_all_variables()` (points to the `init` variable)
- `tf.compat.v1.initialize_all_variables()` (points to the `init` variable)
- `tf.compat.v1.Session()` (points to the `sess1` and `sess2` variables)
- `tf.compat.v1.assign_add()` (points to the `sess1.run(my_var.assign_add(5))` and `sess2.run(my_var.assign_add(2))` lines)

# TensorFlow **changing** variables

- If you'd like to reset your Variables to their starting value, simply call **tf.global\_variables\_initializer()** again
- Or you can use **tf.variables\_initializer()** if you only want to reset a subset of them (see previous slides):

```
# Create Ops
my_var = tf.Variable(0)
init = tf.initialize_all_variables()

# Start Session
sess = tf.Session()

# Initialize Variables
sess.run(init)

# Change the Variable
sess.run(my_var.assign(10))

# Reset the Variable to 0, its initial value
sess.run(init)
```

The diagram illustrates updates to the TensorFlow code. A blue arrow points from `tf.initialize_all_variables()` to a callout box containing `tf.compat.v1.initialize_all_variables()`. Another blue arrow points from `tf.Session()` to a callout box containing `tf.compat.v1.Session()`.

# TensorFlow trainable variables

- **Optimizer** classes automatically train machine learning models
- If there are **Variables** in your graph that should only be changed manually and not with an **Optimizer**, you need to set their trainable parameter to False when creating them:

```
not_trainable = tf.Variable(0, trainable=False)
```